

**The University of Alabama in Huntsville**  
**ECE Department**  
**CPE 431 01**  
**Test 2 Solution**  
**Fall 2018**

1. (3 points) The three C's of cache misses are capacity, conflict, and compulsory.
2. (1 point) Virtual memory is a technique that uses main memory as a "cache" for secondary storage.
3. (1 point) A context switch is a changing of the internal state of the processor to allow a different process to use the processor that includes saving the state needed to return to the currently executing process.
4. (10 points) Calculate the total number of bits required for a cache with the parameters listed below, assuming a 32-bit address and byte addressability. Include valid and dirty bits.

**Cache Data Size:** 32 KiB

**Cache Block Size:** 2 64-bit words

**Cache Set Associativity:** 4-way

**Cache Access Time:** 1 cycle

**32 KiB x 1 word/8 bytes x 1 block/2 words x 1 set/4 blocks = 512 sets**

**Byte offset = 3 bits (8 bytes per word)**

**Block offset = 1 bit (2 words per block)**

**Index = 9 bits (512 sets)**

**Tag = 32 – (9 + 1 + 3) = 32 – 13 = 19 bits**

**Each block has data plus tag plus dirty and valid bits**

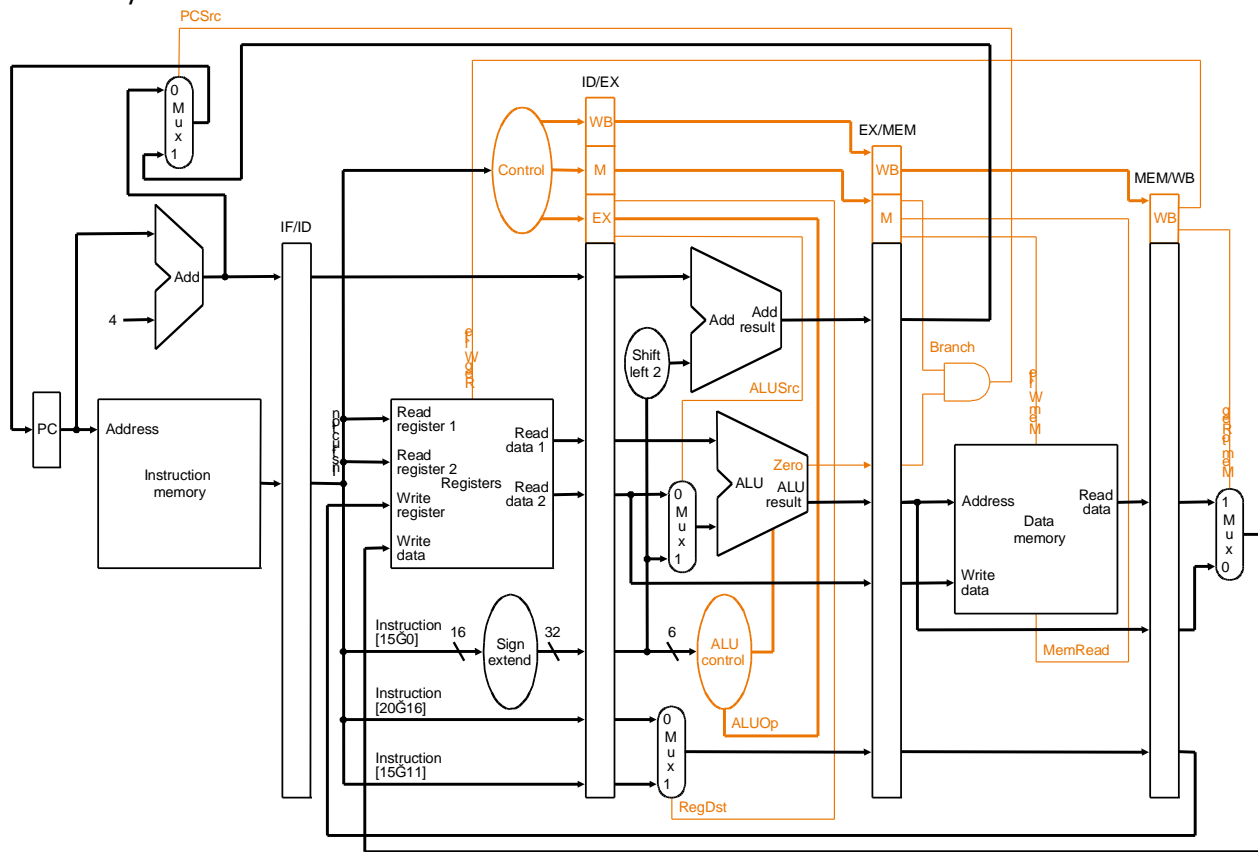
**Total Bits = 512 sets x 4 blocks/set x (128 + 19 + 2) bits/block = 2048 x 149 bits = 305,152 bits**

5. (20 points) Here is a series of address references given as byte addresses: 118, 483, 2069, 321, 368, 1505, 812, 2832, 373, 1411, 511, 122, 690, 4820, 1714, 1508, 2070, 1080. Assuming a two-way set associative-mapped cache with 4-word blocks and a total size of 16 64-bit words that is initially empty and uses LRU, (a) label each reference in the list as a hit or a miss and (b) show the entire history of the cache, including tag and data. **16 words x 1 block/4 words x 1 set/2 blocks = 2 sets**

Byte Address (Decimal)	Byte Address (Hexadecimal)	Byte Address (Binary)	Offsets	Hit/Miss
118	76	0000 0000 01 1 10 110	22	Miss
483	1E3	0000 0001 11 1 00 011	3	Miss
2069	815	0000 1000 00 0 10 101	21	Miss
321	141	0000 0001 01 0 00 001	1	Miss
368	170	0000 0001 01 1 10 000	16	Miss
1505	5E1	0000 0101 11 1 00 001	1	Miss
812	32C	0000 0011 00 1 01 100	12	Miss
2832	B10	0000 1011 00 0 10 000	16	Miss
373	175	0000 0001 01 1 10 101	21	Miss
1411	583	0000 0101 10 0 00 011	3	Miss
511	1FF	0000 0001 11 1 11 111	31	Miss
122	7A	0000 0000 01 1 11 010	28	Miss
690	2B2	0000 0010 10 1 10 010	18	Miss
4820	12D4	0001 0010 11 0 10 100	20	Miss
1714	6B2	0000 0110 10 1 10 010	18	Miss
1508	5E4	0000 0101 11 1 00 100	4	Miss
2070	816	0000 1000 00 0 10 110	22	Miss
1080	438	0000 0100 00 1 11 000	24	Miss

Set	Tag	Data	Tag	Data
0	5, 22, 32	M[320..351], M[1408..1439], M[2048..2069]	32, 44, 75	M[2048..2079], M[2816..2847], M[4800..4831]
1	1, 5, 12, 7, 10, 23	M[96..127], M[352..383], M[800..831], M[480..511], M[672..703], M[1504..1535]	7, 23, 5, 1, 26, 16	M[480..511], M[1504..1535], M[352..383], M[96..127], M[1696..1727], M[1056..1087]

6. (15 points) Consider executing the following code on a pipelined datapath like the one shown except that 1) it supports  $j$  instructions that complete in the ID stage, and 2) it has full forwarding. The register file supports writing in the first half cycle and reading in the second half cycle.



sort:	addi	\$sp, \$sp, -20	208	lw	\$t4, 4(\$t2)	
	sw	\$ra, 16(\$sp)	212	slt	\$t0, \$t4, \$t3	
	sw	\$s3, 12(\$sp)	216	beq	\$t0, \$zero, exit2	
	sw	\$s2, 8(\$sp)	220	add	\$a0, \$s2, \$zero	
	sw	\$s1, 4(\$sp)	224	add	\$a1, \$s1, \$zero	
	sw	\$s0, 0(\$sp)	228	jal	swap	
	add	\$s2, \$a0, \$zero	232	addi	\$s1, \$s1, -1	
	add	\$s3, \$a1, \$zero	236	j	for2tst	
	add	\$s0, \$zero, \$zero	240	exit2:	addi	\$s0, \$s0, 1
for1tst:	slt	\$t0, \$s0, \$s3	244	j	for1tst	
	beq	\$t0, \$zero, exit1	248	exit1:	lw	\$s0, 0(\$sp)
	addi	\$s1, \$s0, -1	252	lw	\$s1, 4(\$sp)	
for2tst:	slt	\$t0, \$s1, \$zero	256	lw	\$s2, 8(\$sp)	
	bne	\$t0, \$zero, exit2	260	lw	\$s3, 12(\$sp)	
	add	\$t1, \$s1, \$s1	264	lw	\$ra, 16(\$sp)	
	add	\$t1, \$t1, \$t1	268	addi	\$sp, \$sp, 20	
200	add	\$t2, \$s2, \$t1	272	jr	\$ra	
204	lw	\$t3, 0(\$t2)				

If the `add $t2` instruction four instructions after the `for2tst` label begins executing in cycle 1 and the `beq $t0, $zero, exit2` is taken, what instructions are found in each of the five stages of the pipeline in the 11<sup>th</sup> cycle? Show the instructions being executed in each stage of the pipeline during each cycle. What value is stored in the ReadData1 field of the ID/EX pipeline register in the 11<sup>th</sup> cycle? Assume that before the instructions are executed, the state of the machine was as follows:

The PC has the value 200<sub>10</sub>, the address of the `add $t2` instruction

Every register has the initial value 20<sub>10</sub> plus the register number.

Every memory word accessed as data has the initial value 10000<sub>10</sub> plus the byte address of the word.

Cycle	IF	ID	EX	MEM	WB
1	<code>add \$t2</code>				
2	<code>lw \$t3</code>	<code>add \$t2</code>			
3	<code>lw \$t4</code>	<code>lw \$t3</code>	<code>add \$t2</code>		
4	<code>slt \$t0</code>	<code>lw \$t4</code>	<code>lw \$t3</code>	<code>add \$t2</code>	
5	<code>beq \$t0</code>	<code>slt \$t0</code>	<code>lw \$t4</code>	<code>lw \$t3</code>	<code>add \$t2</code>
6	<code>beq \$t0</code>	<code>slt \$t0</code>	bubble	<code>lw \$t4</code>	<code>lw \$t3</code>
7	<code>add \$a0</code>	<code>beq \$t0</code>	<code>slt \$t0</code>	bubble	<code>lw \$t4</code>
8	<code>add \$a1</code>	<code>add \$a0</code>	<code>beq \$t0</code>	<code>slt \$t0</code>	bubble
9	<code>jal swap</code>	<code>add \$a1</code>	<code>add \$a0</code>	<code>beq \$t0</code>	<code>slt \$t0</code>
10	<code>addi \$s0</code>	bubble	bubble	bubble	<code>beq \$t0</code>
11	<code>j for1test</code>	<code>addi \$s0</code>	bubble	bubble	bubble

Instruction of interest is the bubbled `jal swap`. Assuming `swap` is close by, the `rs` field is 0, so the value is 20.

ID/EX.ReadData1 = 20

7. (20 points) Consider the execution of the following loop in a statically scheduled superscalar processor that has full forwarding. Additionally, any branches are resolved in the EX stage.

```

Loop:  lw    $1, 40($6)
       add   $1, $7, $1
       sw    $1, 20($5)
       addi  $6, $6, 4
       addi  $5, $5, -4
       bne   $5, $0, Loop

```

- (15 points) Unroll this loop so that four iterations of it are done at once and schedule it for maximum performance on a 2-issue static superscalar processor that has one slot for R-type/branch instructions and one slot for lw/sw instructions. Assume that the loop always executes a number of iterations that is a multiple of 4. You can use registers **\$10** through **\$20** when changing the code to eliminate dependences.
- (5 points) Calculate the number of cycles for the original and for the unrolled, scheduled code and the speedup of unrolled, scheduled compared to the original.

a. Unrolled, scheduled

Cycle	R-type/branch	lw/sw
1	addi \$5, \$5, -16	lw \$1, 40(\$6)
2		lw \$10, 440(\$6)
3	add \$1, \$7, \$1	lw \$11, 48(\$6)
4	add \$10, \$7, \$10	lw \$12, 52(\$6)
5	add \$11, \$7, \$11	sw \$1, 36(\$5)
6	add \$12, \$7, \$12	sw \$10, 32(\$5)
7	addi \$6, \$6, 16	sw \$11, 28(\$5)
8	bne \$5, \$0, Loop	sw \$12, 24(\$5)

b. Unrolled, scheduled ET = 8 cycles (issue slots) + 2 cycles control stall  
 Original ET = 4 \* (6 cycles (issue slots) + 1 data hazard stall (lw-add) + 2 cycles control stall)  
 Speedup = Original ET/Unrolled, scheduled ET = 36/10 = 3.6

8. (15 points) Virtual memory uses a page table to track the mapping of virtual addresses to physical addresses. The following table is a stream of virtual addresses as seen on a system. Assume 32 KiB pages, byte addressing, a four-entry fully associative TLB, and true LRU replacement. If pages must be brought in from disk, increment the next largest page number. Given the address stream, and the shown initial state of the TLB and page table, show the final state of the system. Also list for each reference if it is a hit in the page table, or a page fault.

#### TLB

Valid	Tag	Physical Page Number
1	<del>11</del> , 1	<del>12</del> , 13
1	7	4
1	3	6
0, 1	4, 5, 4	9, <del>11</del> , 9

#### Page table

VPN	Valid	Physical page or in disk
0	1	5
1	0	Disk, <del>13</del>
2	0	Disk
3	1	6
4	1	9
5	1	11
6	0	Disk
7	1	4
8	0	Disk
9	0	Disk
10	1	3
11	1	12
12	0	Disk

Address	VPN	TLB Hit/Miss	Page Table Hit/Miss	Page Fault Y/N
184,760	5	M	H	N
110,824	3	H	-	-
236,100	7	H	-	-
119,200	3	H	-	-
56,312	1	M	M	Y
151,692	4	M	H	N
126,856	3	H	-	-
40,000	1	H	-	-

9. (10 points) Consider a pipeline for a register-memory architecture. The architecture has two instruction formats: a register-register format and a register-memory format. There is a single memory addressing mode (offset + base register). There is a set of ALU operations as follows:

$$\text{Rdest} \leftarrow \text{Rsrc1 ALUop Rsrc2}$$

**or**       $R_{dest} \leftarrow R_{src1} \text{ ALUop } MEM[R_{src2} + \text{offset}]$

or  $Rdest \leftarrow MEM[Rsrc2 + offset]$

**or**      $\text{MEM}[\text{Rsrc2} + \text{offset}] \leftarrow \text{Rsrc1}$

where the ALUop is one of the following:

Add, Subtract, And, Or (with or without offset), Load (Rsrc1 omitted), Store (Rdest omitted)

Rdest, Rsrc1 and Rsrc2 are registers.

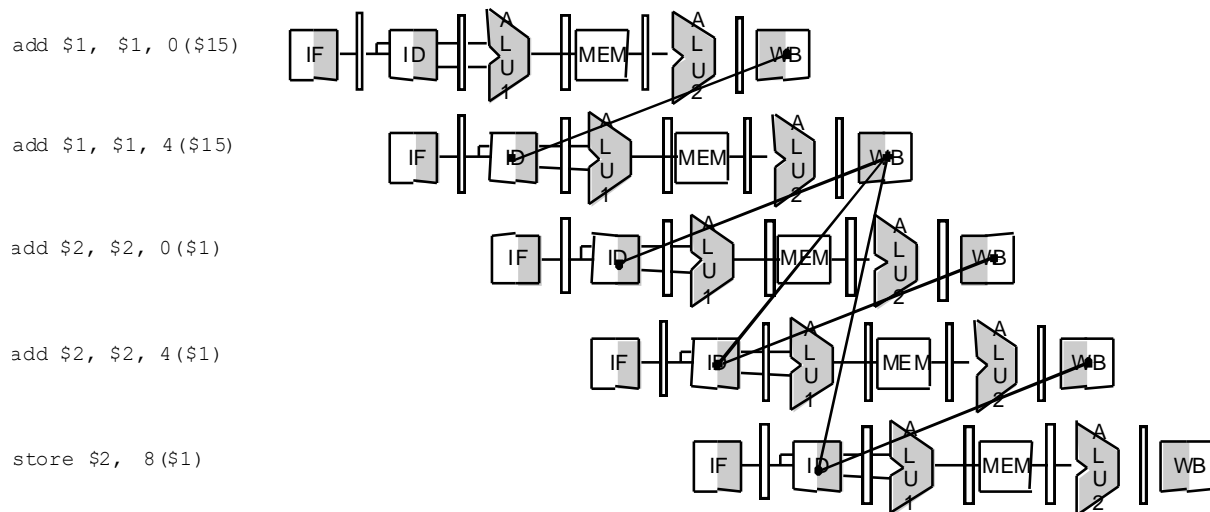
Branches use a full compare of two registers and are PC-relative. Assume that this machine is pipelined so that a new instruction is started every clock cycle. The pipeline structure is

IF	RF	ALU1	MEM	ALU2	WB						
	IF	RF	ALU1	MEM	ALU2	WB					
		IF	RF	ALU1	MEM	ALU2	WB				
			IF	RF	ALU1	MEM	ALU2	WB			
				IF	RF	ALU1	MEM	ALU2	WB		
					IF	RF	ALU1	MEM	ALU2	WB	
						IF	RF	ALU1	MEM	ALU2	WB

The first ALU stage is used for effective address calculation for memory references and branches. The second ALU stage is used for operations and branch comparisons. RF is both a decode and register-fetch stage. Assume that when a register read and a register write of the same register occur in the same clock cycle, the write data is forwarded. For the following code fragment:

```
add    $1, $1, 0($15)
add    $1, $1, 4($15)
load   $2, 0($1)
add    $2, $2, 4($1)
store  $2, 8($1)
```

identify the data dependencies and draw a figure (using the multiple clock style) showing them as lines. You have space on the next page for this figure if you choose to use it.



10. (5 points) The following code is written in C, where elements within the same row are stored contiguously. Does  $B(j, 0)$  exhibit spatial locality? temporal locality? Explain your answers. A, B, and C are all arrays of integers 8000 by 8000.

```
for (i = 0; i < 8000; i++)
    for (j = 0; j < 8000; j++)
        A(i, j) = B(j, 0) + A(j, i) + C(0, i);
```

$B(j, 0)$  does not exhibit spatial locality, the access is  $B(0, 0)$ ,  $B(1, 0)$ , ... . Since they don't come from the same row, they are stored 8000 locations apart.

$B(j, 0)$  exhibits temporal locality if once every 8000 times is soon, you could also argue that it doesn't because it changes every time in the inner loop